


Resource Consumption & Latency Profile Across Model Architectures

Energy Consumption per Training Run (mWh)



Inference Latency per 1,000 Samples (ms)

